

AQI PREDICTOR

Serverless Air Quality Forecasting for Islamabad & Rawalpindi

**A Machine Learning System for 72-Hour Advanced AQI
Prediction**

Author

Zeeshan Ali Syed

Mentors

Mam Hafsa & Mam Umema

**10 Pearls Shine Internship Program | Data Sciences Track,
Cohort 9**

1. EXECUTIVE SUMMARY

Air quality forecasting is becoming critical for public health in Islamabad and Rawalpindi, where seasonal smog poses serious respiratory risks. This project presents **AQI Predictor**, a fully automated, serverless machine learning system that forecasts air quality up to 72 hours in advance.

Key Achievements:

- Validation R^2 Scores: 0.84 (24h), 0.78 (48h), 0.70 (72h) using XGBoost
- Zero Operational Cost: Deployed entirely on free cloud platforms
- Production-Ready: 24/7 automated pipelines with no manual intervention
- Interpretable: SHAP values explain which factors drove each prediction

The system ingests real-time data hourly, trains models daily, and serves predictions through an interactive web dashboard. All components operate autonomously through GitHub Actions orchestration.

2. PROBLEM STATEMENT & MOTIVATION

The Air Quality Crisis in Islamabad & Rawalpindi

Every winter (November–February), AQI levels in Islamabad and Rawalpindi regularly exceed 200, classified as "Very Unhealthy" by the EPA. Even healthy people experience respiratory discomfort; vulnerable populations face serious health risks. This is caused by agricultural crop burning in Punjab combined with meteorological factors like temperature inversions and stagnant winds that trap pollutants near ground level.

The Core Problem:

Air quality reporting is reactive, not predictive. By the time someone checks the AQI and sees dangerously high levels, they're already breathing hazardous air. There's no time to adjust plans, prepare protection, or take preventive action.

Why This Matters: Air quality follows predictable meteorological patterns. Machine learning can learn these patterns and predict AQI 24–72 hours ahead, enabling informed decisions, health protection, and evidence-based policy interventions.

3. PROJECT OBJECTIVES & APPROACH

Project Goals:

- Build a fully automated ML pipeline with zero manual intervention
- Achieve interpretable predictions (SHAP values explain each forecast)
- Minimize operational overhead (free cloud platforms only)
- Enable 72-hour forecasting (24h, 48h, and 72h separate models)
- Serve predictions publicly via interactive web dashboard

Why We Chose This Approach:

Serverless Architecture , No servers to manage. Use managed services (Hopworks, GitHub Actions, Streamlit) that handle infrastructure automatically. Cost-effective and scalable.

Multiple Models for Horizons , One model predicting all three timeframes is convenient but suboptimal. Separate models for 24h, 48h, 72h allow each to specialize. A 72-hour forecast has inherently lower accuracy, separating models prevents this degradation from affecting 24-hour predictions.

Daily Retraining — Air quality patterns shift with seasons and long-term trends. Retraining daily keeps models calibrated. Too frequent risks overfitting to noise; too infrequent risks staleness.

4. SYSTEM ARCHITECTURE & DESIGN

Four integrated components orchestrated through GitHub Actions:

4.1 Feature Pipeline (Runs Hourly)

Fetches current AQI readings (AQICN) and weather data (Open-Meteo) every hour. Engineers 35 numerical features from raw data. Persists features to Hopsworks feature store (centralized data warehouse). Idempotent, running multiple times on the same timestamp produces no duplicates.

Why Hourly? Both data sources update hourly. Hourly runs ensure we always have the latest data. More frequent runs would waste API quotas without improving forecasts.

4.2 Training Pipeline (Runs Daily at 7:00 AM PKT)

Pulls all historical features from Hopsworks. Computes derived features (lags, rolling averages). Splits chronologically into training and validation (no shuffling). Trains three algorithms (Ridge, Random Forest, XGBoost) for each forecast horizon (24h, 48h, 72h). Selects best performer (lowest validation RMSE). Computes SHAP feature importance. Registers model in Hopsworks registry with version metadata.

Why Daily? Air quality patterns shift with seasons. Daily retraining calibrates models to recent data distribution. Failed or underperforming models never get deployed, system automatically falls back to previous best version.

4.3 Model Registry (Hopsworks)

Maintains versioned models. Dashboard always queries the version with lowest validation RMSE for each horizon. Automatic versioning enables rollback if a model degrades.

4.4 Streamlit Dashboard

Web interface deployed on Streamlit Community Cloud. Queries feature store for latest observations. Loads best model from registry. Generates forecasts with SHAP explanations. Renders interactive visualizations. Displays health advisories based on AQI levels.

Live at: <https://10pearlsss-zeeshan.streamlit.app/>

5. DATA SOURCES & FEATURE ENGINEERING

5.1 AQICN (Air Quality Index China)

Ground-level sensor measurements from Islamabad station @517168. Provides hourly readings of PM2.5, PM10, NO2, CO, O3, SO2. This is ground truth, physical sensors on the ground, not atmospheric models. Free API token (500 requests/month); we use 24 per day.

5.2 Open-Meteo

Completely free weather API requiring no authentication. Provides hourly weather (temperature, wind, humidity, pressure, precipitation). Historical data back to 1940 for backfilling.

5.3 Feature Engineering: 35 Features from Raw Data

Pollutants (7): aqi, pm25, pm10, no2, co, o3, so2 , Direct sensor readings

Weather (9): temperature, humidity, wind_speed, pressure, precipitation, wind_u, wind_v , Atmospheric conditions; wind decomposed into u,v to avoid circular discontinuity where $359^\circ \neq 1^\circ$ numerically but $\approx 1^\circ$ physically

Temporal (9): hour_sin, hour_cos, month_sin, month_cos, is_rush_hour, is_weekend , Diurnal and seasonal cycles; sine/cosine encoding handles time circularity

Lags & Rolling (9): aqi_lag_1h/3h/6h/24h, aqi_change, rolling averages (3h, 6h, 24h), rolling_std , Autoregressive persistence and momentum

Why These Features?

Lag Features: AQI has strong autoregressive behavior. The 24-hour lag (AQI same hour yesterday) is often the single strongest predictor, pollution patterns repeat daily.

Rolling Averages: Capture momentum. Rapidly changing AQI signals different atmospheric conditions than stable AQI.

Wind Decomposition: Wind direction (0-360°) is circular. 359° is close to 1° physically but distant numerically. Decomposing into u,v components (east-west, north-south) removes this artifact.

6. MODEL DEVELOPMENT & ALGORITHM SELECTION

6.1 Chronological Train-Test Split (No Shuffling)

Standard practice in ML is to shuffle data before splitting. But in time-series forecasting, shuffling introduces information leakage. If your validation set contains data from every timestamp in history, your model has "seen" future data, validation metrics become too optimistic.

Our Approach: Strict chronological split. Training = all data except final 15%. Validation = most recent 15% (~26 days). This mimics real-world deployment where models must predict timestamps they've never seen.

6.2 Three Algorithm Comparison

Ridge Regression: Strong linear baseline. Captures additive relationships but lacks nonlinear flexibility for complex atmospheric interactions.

Random Forest: Ensemble of decision trees. Captures nonlinear relationships and feature interactions automatically. However, showed overfitting on validation sets—model was memorizing rather than generalizing.

XGBoost: Extreme Gradient Boosting. Builds shallow trees sequentially, each correcting mistakes of previous ones. Key advantage: built-in L1/L2 regularization prevents overfitting. Learning rate control ensures stable, predictable training.

Results—XGBoost Won:

- 24-hour: RMSE 12.9 (22% better than RF), R^2 0.84
- 48-hour: RMSE 16.4 (15% better than RF), R^2 0.78
- 72-hour: RMSE 20.2 (13% better than RF), R^2 0.70

Why XGBoost? Superior regularization prevents overfitting. Learning rate control ensures stable training. 22% RMSE reduction at 24-hour horizon is significant, maintains nonlinear expressiveness while avoiding memorization.

6.3 SHAP Feature Importance: What Drives Predictions?

After each daily training run, we compute SHAP (SHapley Additive exPlanations) values, a game-theoretic approach to feature importance. SHAP allocates credit fairly to each feature based on its contribution to individual predictions.

Top Features:

1. aqi_roll_24h (14.2) -Daily persistence is dominant. AQI is highly autocorrelated.
2. aqi_lag_24h (11.8) - Same hour yesterday is the single strongest lag predictor.
3. pm25 (8.9) -Direct PM2.5 concentration drives AQI calculation.
4. wind_speed (7.3) - Wind disperses pollutants (primary atmospheric mechanism).

5. `aqi_lag_6h` (6.1) - Recent trends signal momentum shifts in pollution.

Why This Matters: Instead of a black box that says "AQI will be 180 tomorrow," we can explain: "AQI will be 180 because today's AQI is high (persists), wind is weak (can't disperse), and it's winter (seasonal effect)." Stakeholders trust what they understand.

7. AUTOMATION WITH GITHUB ACTIONS

Once deployed, the system runs 24/7 with zero human intervention. GitHub Actions triggers two pipelines on schedule:

- **Feature Pipeline:** Every hour (2–3 min runtime, 720 times/month)
- **Training Pipeline:** Daily at 7:00 AM PKT (12 min runtime, 30 times/month)

GitHub Quota Management:

Free tier provides 2000 minutes/month for public repos. Our usage:

- Feature: $720 \times 2.5 \text{ min} = 1800 \text{ min/month}$
- Training: $30 \times 12 \text{ min} = 360 \text{ min/month}$
- **Total:** 2160 min/month (exceeds quota by 160 min initially)

Solution: Optimized training from 18 min to 12 min through vectorized feature computation and efficient serialization. Now run safely under quota with buffer for growth.

8. DEPLOYMENT ARCHITECTURE (ZERO COST)

Infrastructure Stack-All Free Tiers:

- **Hopsworks:** Feature store + model registry. Unlimited read/write. Free tier.
- **GitHub Actions:** Orchestration. 2000 min/month included with public repos.
- **Streamlit Community Cloud:** Web frontend. Free hosting for one active project.
- **Open-Meteo & AQICN:** Data APIs. No infrastructure cost.

Total Monthly Cost: \$0

Why This Matters:

Free tiers aren't temporary offers. These are permanent, stable services from major companies. By designing within free quotas from day one, the system is sustainable long-term. No surprise bills, no need for funding to keep systems running.

9. CHALLENGES & SOLUTIONS

Challenge 1: AQICN Sensor Failures

Sensors occasionally go offline or return incomplete data.

Solution: Implemented safe extraction. Missing values become NaN instead of crashing. Model handles NaN through imputation. Pipeline completes gracefully even when sensors fail.

Challenge 2: Wind Direction Circularity

Wind direction (0-360°) is circular: 359° and 1° are physically nearly identical but numerically distant. Linear models treat them as far apart, introducing error.

Solution: Decomposed wind into orthogonal *u* (east-west) and *v* (north-south) components. Circular continuity is automatic.

Challenge 3: GitHub Quota Management

Initial design used 2250 minutes (over 2000-minute quota).

Solution: Optimized training pipeline from 18 to 12 minutes through vectorized features and efficient serialization. Now safely under quota.

Challenge 4: Feature Store Availability

If Hopworks goes down, dashboard crashes.

Solution: Implemented fallback caching. Dashboard attempts Hopworks first, but falls back to cached model stored in GitHub if unavailable.

10. KEY LESSONS & FUTURE WORK

Lessons Learned:

1. Autoregressive features are powerful. The 24-hour lag explains 50%+ of AQI variance alone.
2. Data quality > algorithm sophistication. Ensuring low NaN rates and valid ranges had bigger impact than ensemble tuning.
3. Chronological splits are non-negotiable. Discipline to hold out only future data matters more than maximizing training size.
4. Interpretability builds trust. SHAP values converted abstract weights into actionable insights. Stakeholders trust what they understand.
5. Respect free tier constraints. Designing within quotas from day one ensures long-term sustainability.

Future Roadmap:

- Probabilistic Forecasting: Quantile regression for confidence intervals (10th, 50th, 90th percentiles)
 - Sequence Models: LSTM/GRU for multi-step temporal dependencies. Ensemble with XGBoost.
 - Multi-City Expansion: Adapt to Lahore, Karachi, Peshawar with city-specific AQICN IDs
 - Data Fusion: Integrate NASA MODIS satellite and EPA-AirNow data for 5–10% RMSE improvement
-

11. CONCLUSION

This project proves that production-grade air quality forecasting is achievable with minimal infrastructure and zero operational cost. Using serverless platforms, open-source ML libraries, and free data APIs, we built a system that predicts AQI 72 hours ahead with R^2 scores exceeding 0.70.

Impact:

For Residents: Advance warning enables informed decisions about outdoor activities, respiratory protection, and personal health.

For Policymakers: Daily retraining and SHAP rankings reveal which meteorological factors most influence air quality, informing evidence-based environmental policy.

Sustainability:

The modular, serverless architecture is replicable to other cities and scalable to new data sources. As air quality forecasting becomes critical for public health, this open-source template can serve other regions.

GitHub Repository: <https://github.com/ZeeshanAli77/AQIPredictor-10pearls>

Live Dashboard: <https://10pearlsss-zeeshan.streamlit.app/>

Contact: zeeshann.2003@gmail.com